


Case Study: Diversity and Innovation

Planning a model: the response would be novelty, and predictors are sex and race ethnicity. I wanted to know if sex and race would have any correlation with innovation level.


```
innovation <- read_csv('https://sldr.netlify.app/data/phd_innovation.csv', show_col_types = TRUE)
model <- lm(novelty ~ sex + race_ethnicity, data = innovation)
nrow(innovation)
```

```
[1] 4195
```


```
gf_boxplot(novelty ~ sex, data = innovation)
```



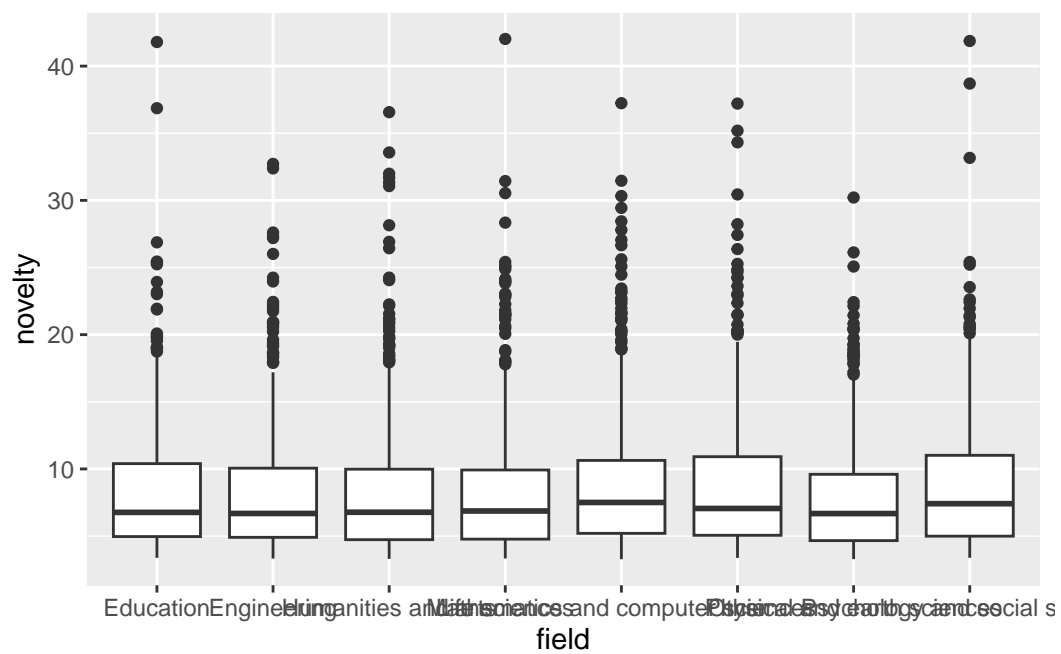
```
gf_boxplot(novelty ~ race_ethnicity, data = innovation)
```



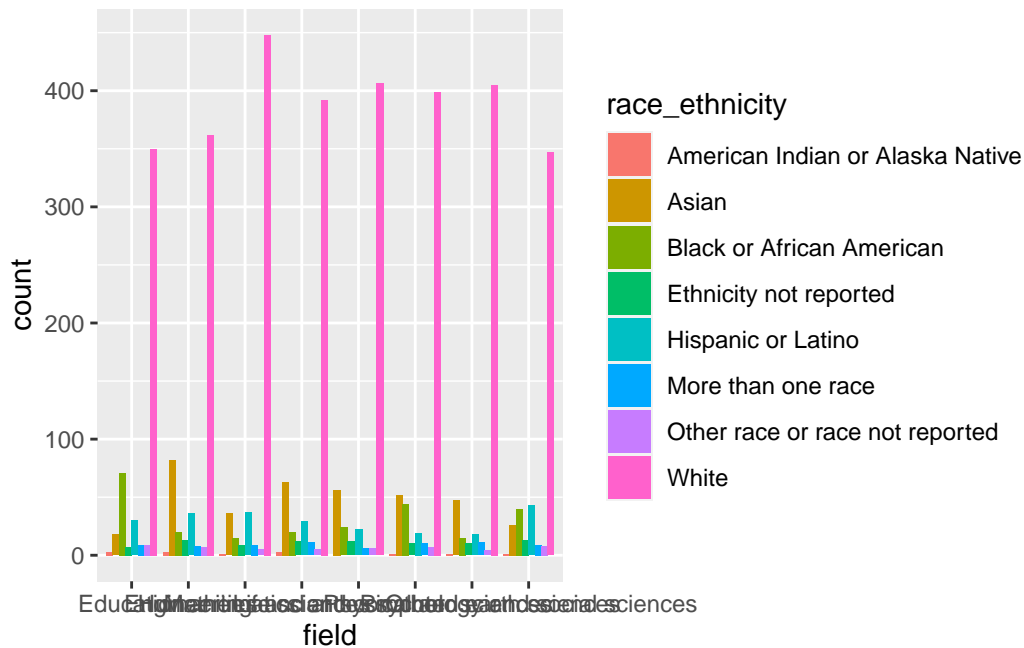
```
gf_point(novelty ~ year, data = innovation)
```



```
gf_boxplot(novelty ~ field, data = innovation)
```



```
gf_bar( ~ field,
        fill = ~ race_ethnicity,
        data = innovation,
        position = 'dodge')
```



The graph between novelty and sex doesn't show much since they have very similar median and IQR. The graph between novelty and race ethnicity shows a dominant lead for White followed by Asians and Black or African American. I can also see a small trend as the year goes on the novelty level tend to be down a little bit. There are some fields are leading the novelty level but I can't really tell since the names are all overlapped.

```
summary(model)
```

Call:

```
lm(formula = novelty ~ sex + race_ethnicity, data = innovation)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.313	-3.551	-1.491	1.870	33.519

Coefficients:

	Estimate	Std. Error	t value
(Intercept)	8.03157	1.38423	5.802
sexMale	-0.09716	0.15580	-0.624
race_ethnicityAsian	0.85802	1.40703	0.610
race_ethnicityBlack or African American	0.80241	1.41871	0.566
race_ethnicityEthnicity not reported	1.22550	1.48426	0.826
race_ethnicityHispanic or Latino	0.83296	1.42123	0.586
race_ethnicityMore than one race	1.16157	1.50115	0.774
race_ethnicityOther race or race not reported	0.72359	1.54965	0.467
race_ethnicityWhite	0.34103	1.38608	0.246

	Pr(> t)
(Intercept)	7.03e-09 ***
sexMale	0.533
race_ethnicityAsian	0.542
race_ethnicityBlack or African American	0.572
race_ethnicityEthnicity not reported	0.409
race_ethnicityHispanic or Latino	0.558
race_ethnicityMore than one race	0.439
race_ethnicityOther race or race not reported	0.641
race_ethnicityWhite	0.806

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.986 on 4186 degrees of freedom

Multiple R-squared: 0.002523, Adjusted R-squared: 0.0006166

F-statistic: 1.323 on 8 and 4186 DF, p-value: 0.2265

Causal Diagram:

My main predictor is race ethnicity with my response variable novelty. I made sex as a moderator since it doesn't affect novelty level directly. Both field and year showed a reasonable correlation with novelty. I would include race ethnicity, field, and year as my predictors to my new model.

```
nmodel <- lm(novelty ~ race_ethnicity + field + year, data = innovation)
summary(nmodel)
```

Call:

```
lm(formula = novelty ~ race_ethnicity + field + year, data = innovation)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.656	-3.522	-1.545	1.876	33.635

Coefficients:

	Estimate	Std. Error	t value
(Intercept)	-30.53964	27.54555	-1.109
race_ethnicityAsian	0.70139	1.40565	0.499
race_ethnicityBlack or African American	0.56545	1.41834	0.399
race_ethnicityEthnicity not reported	1.03825	1.48281	0.700
race_ethnicityHispanic or Latino	0.61522	1.42053	0.433
race_ethnicityMore than one race	0.99800	1.50053	0.665
race_ethnicityOther race or race not reported	0.50596	1.54801	0.327
race_ethnicityWhite	0.18061	1.38506	0.130
fieldEngineering	-0.08386	0.31417	-0.267
fieldHumanities and arts	0.09449	0.30943	0.305
fieldLife sciences	-0.03971	0.31287	-0.127
fieldMathematics and computer sciences	0.56017	0.31306	1.789
fieldOther	0.44429	0.31044	1.431
fieldPhysical and earth sciences	-0.51178	0.31651	-1.617
fieldPsychology and social sciences	0.45058	0.31835	1.415
year	0.01920	0.01371	1.400

Pr(>|t|)

(Intercept)	0.2676
race_ethnicityAsian	0.6178
race_ethnicityBlack or African American	0.6902
race_ethnicityEthnicity not reported	0.4838
race_ethnicityHispanic or Latino	0.6650
race_ethnicityMore than one race	0.5060
race_ethnicityOther race or race not reported	0.7438
race_ethnicityWhite	0.8963
fieldEngineering	0.7895
fieldHumanities and arts	0.7601
fieldLife sciences	0.8990
fieldMathematics and computer sciences	0.0736 .
fieldOther	0.1525
fieldPhysical and earth sciences	0.1060
fieldPsychology and social sciences	0.1570
year	0.1615

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.978 on 4179 degrees of freedom


Multiple R-squared: 0.007298, Adjusted R-squared: 0.003735
 F-statistic: 2.048 on 15 and 4179 DF, p-value: 0.009758

$$\begin{aligned}
 y = & -30.54 + 0.7\beta_{\text{raceAsian}} + 0.57\beta_{\text{raceAfricanAmerican}} + 1.04\beta_{\text{raceNotReported}} \\
 & + 0.62\beta_{\text{raceLatino}} + 1\beta_{\text{raceMorethanone}} + 0.51\beta_{\text{raceOther}} + 0.18\beta_{\text{raceWhite}} \\
 & - 0.08\beta_{\text{fieldEngineering}} + 0.09\beta_{\text{fieldHumanityandArts}} - 0.04\beta_{\text{fieldLifeScience}} \\
 + & 0.56\beta_{\text{fieldMathandCS}} + 0.44\beta_{\text{fieldOther}} - 0.51\beta_{\text{fieldPhysicalandEarthScience}} \\
 & + 0.45\beta_{\text{fieldPhychologyandSocialScience}} + 0.02\beta_{\text{year}} + \epsilon, \\
 & \epsilon \sim N(0, 4.978)
 \end{aligned}$$

```

innovation$residuals <- residuals(nmodel)
gf_point(residuals ~ fitted(nmodel), data = innovation)


```




```

gf_histogram(~residuals, data = innovation, bins=100)

```



```
s245::gf_acf(~nmodel) |> gf_lims(y = c(-1,1))
```



Residuals and fitted model: It doesn't really show linearity from this graph, so I don't think it meets the condition.


Histogram of residuals: This is not normally distributed. This is a right-skewed graph. It also don't meet the condition.

ACF plot of residuals: There is only one significant spike, so I would say this meets the condition.

Since not all of them met the conditions, I cannot draw a valid conclusion from this model.

```
b2 <- innovation |>
  select(field, race_ethnicity, year, novelty) |>
  na.omit() |>
  mutate(preds = predict(nmodel),
         resids = resid(nmodel))

gf_point(novelty ~ race_ethnicity,
  data = b2) |>
gf_line(preds ~ race_ethnicity)
```



I don't know why somehow the line doesn't show.